

```
• (base) kindl@Kindle:~/llm-challenge$ curl -sS --max-time 120 http://127.0.0.1:30000/v1/chat/completions -H 'Content-Type: application/json' --data-raw '{"model": "Qwen/Qwen3-0.6B", "messages": [{"role": "user", "content": "请用一句话介绍 SGLang。"}], "max_tokens": 32}'
{"id": "a2eeb146eba94321b7f600c9f867a750", "object": "chat.completion", "created": 1790588825, "model": "Qwen/Qwen3-0.6B", "choices": [{"index": 0, "message": {"role": "assistant", "content": "<think>\n好的，用户让我用一句话介绍SGLang。首先，我需要确认SGLang是什么。SGLang是Simplified Golang的", "reasoning_content": null, "tool_calls": null}, "logprobs": null, "finish_reason": "length", "matched_stop": null}], "usage": {"prompt_tokens": 16, "total_tokens": 48, "completion_tokens": 32, "prompt_tokens_details": null, "reasoning_tokens": 0}, "metadata": {"weight_version": "default"}}(base) kindl@Kindle:~/llm-challenge$
• llenge$ curl -sS http://127.0.0.1:3000 curl -sS http://127.0.0.1:30000/v1/models
{"object": "list", "data": [{"id": "Qwen/Qwen3-0.6B", "object": "model", "created": 1790589191, "owned_by": "sglang", "root": "Qwen/Qwen3-0.6B", "parent": null, "max_model_len": 4096}]}(base) kindl@Kindle:~/llm-challenge$
```

```

(base) kindl@Kindle:~/llm-challenge$ cd /mnt/c/Users/kindle/Desktop/训练营/挑战一/output/workload/runs/poisson-20261008-pytorch
cat results.jsonl
cat summary.json
{"request_id": 1, "trace_index": 1724, "trace_input_length": 1309, "trace_output_length": 451, "target_input_tokens": 1309, "client_input_tokens": 1309, "client_chat_template_tokens": 1317, "max_output_tokens": 451, "planned_send_seconds": 1.054093147903962, "actual_send_seconds": 1.056017151991854, "dispatch_lateness_seconds": 0.0019240040952235304, "http_status": 200, "status": "ok", "prompt_tokens": 1317, "completion_tokens": 332, "finish_reason": "stop", "latency_seconds": 4.220701393000127, "complete_seconds": 5.27671903800001}
{"request_id": 6, "trace_index": 9101, "trace_input_length": 2013, "trace_output_length": 432, "target_input_tokens": 2013, "client_input_tokens": 2013, "client_chat_template_tokens": 2021, "max_output_tokens": 432, "planned_send_seconds": 6.188226591992954, "actual_send_seconds": 6.18993357600084, "dispatch_lateness_seconds": 0.0017069840078853815, "http_status": 200, "status": "ok", "prompt_tokens": 2021, "completion_tokens": 221, "finish_reason": "stop", "latency_seconds": 1.1935807059999206, "complete_seconds": 7.3835151820003375}
{"request_id": 2, "trace_index": 3495, "trace_input_length": 902, "trace_output_length": 370, "target_input_tokens": 902, "client_input_tokens": 902, "client_chat_template_tokens": 910, "max_output_tokens": 370, "planned_send_seconds": 2.7269513346587066, "actual_send_seconds": 2.7297929859996657, "dispatch_lateness_seconds": 0.0028416513409590394, "http_status": 200, "status": "ok", "prompt_tokens": 910, "completion_tokens": 286, "finish_reason": "stop", "latency_seconds": 2.299317680000968, "complete_seconds": 5.029111165000359}
{"request_id": 3, "trace_index": 6055, "trace_input_length": 1314, "trace_output_length": 342, "target_input_tokens": 1314, "client_input_tokens": 1314, "client_chat_template_tokens": 1322, "max_output_tokens": 342, "planned_send_seconds": 3.9861562159353774, "actual_send_seconds": 3.9880417519998446, "dispatch_lateness_seconds": 0.0018855360644671748, "http_status": 200, "status": "ok", "prompt_tokens": 1322, "completion_tokens": 342, "finish_reason": "length", "latency_seconds": 2.4693911880003725, "complete_seconds": 6.45743347700045}
{"request_id": 4, "trace_index": 8231, "trace_input_length": 2957, "trace_output_length": 14, "target_input_tokens": 2957, "client_input_tokens": 2957, "client_chat_template_tokens": 2965, "max_output_tokens": 14, "planned_send_seconds": 5.0040634810652955, "actual_send_seconds": 5.005740122000134, "dispatch_lateness_seconds": 0.0016766409348383249, "http_status": 200, "status": "ok", "prompt_tokens": 2965, "completion_tokens": 14, "finish_reason": "length", "latency_seconds": 0.2116978529993503, "complete_seconds": 5.217438453000796}
{"request_id": 8, "trace_index": 11009, "trace_input_length": 1704, "trace_output_length": 389, "target_input_tokens": 1704, "client_input_tokens": 1704, "client_chat_template_tokens": 1712, "max_output_tokens": 389, "planned_send_seconds": 7.813531458192071, "actual_send_seconds": 7.814254192999215, "dispatch_lateness_seconds": 0.0007227348071436879, "http_status": 200, "status": "ok", "prompt_tokens": 1712, "completion_tokens": 260, "finish_reason": "stop", "latency_seconds": 1.2858164950011997, "complete_seconds": 9.1000713079999315}
{"request_id": 9, "trace_index": 11663, "trace_input_length": 1581, "trace_output_length": 283, "target_input_tokens": 1581, "client_input_tokens": 1581, "client_chat_template_tokens": 1589, "max_output_tokens": 283, "planned_send_seconds": 9.423614934684801, "actual_send_seconds": 9.425707930000499, "dispatch_lateness_seconds": 0.002092995315697621, "http_status": 200, "status": "ok", "prompt_tokens": 1589, "completion_tokens": 232, "finish_reason": "stop", "latency_seconds": 1.134497682998699, "complete_seconds": 10.5602062880007}
{"request_id": 7, "trace_index": 10578, "trace_input_length": 899, "trace_output_length": 26, "target_input_tokens": 899, "client_input_tokens": 899, "client_chat_template_tokens": 907, "max_output_tokens": 26, "planned_send_seconds": 7.385505156252373, "actual_send_seconds": 7.386938642999667, "dispatch_lateness_seconds": 0.0014334867472935997, "http_status": 200, "status": "ok", "prompt_tokens": 907, "completion_tokens": 26, "finish_reason": "length", "latency_seconds": 0.13433406799958902, "complete_seconds": 7.521273222000673}
{"request_id": 5, "trace_index": 8392, "trace_input_length": 926, "trace_output_length": 138, "target_input_tokens": 926, "client_input_tokens": 926, "client_chat_template_tokens": 934, "max_output_tokens": 138, "planned_send_seconds": 5.512703426431932, "actual_send_seconds": 5.513431109999146, "dispatch_lateness_seconds": 0.0007276835672147541, "http_status": 200, "status": "ok", "prompt_tokens": 934, "completion_tokens": 138, "finish_reason": "length", "latency_seconds": 0.810381719000361, "complete_seconds": 6.32381340000029}
{"request_id": 10, "trace_index": 16324, "trace_input_length": 890, "trace_output_length": 8, "target_input_tokens": 890, "client_input_tokens": 890, "client_chat_template_tokens": 898, "max_output_tokens": 8, "planned_send_seconds": 13.39346158968835, "actual_send_seconds": 13.39783754599921, "dispatch_lateness_seconds": 0.004375956310859408, "http_status": 200, "status": "ok", "prompt_tokens": 898, "completion_tokens": 8, "finish_reason": "length", "latency_seconds": 0.2225833690008585, "complete_seconds": 13.620421399000406}
{
  "count": 10,
  "successes": 10,
  "errors": 0,
  "mean_success_latency_seconds": 1.3982302154001445,
  "max_dispatch_lateness_seconds": 0.004375956310859408,
  "successful_token_counts_present": true,
  "note": "length-filtered synthetic workload, original lengths retained; no cache-performance or TTFT conclusion"
}

```

修正实验原始结果核验

以下由保存的 JSON 数据生成，属于结果附页，不是终端截图。

第1页为单次推理，第2页为新轮次截图；原始数据见下列正式结果目录。

SGLang 0.5.14 | Ray 2.56.0 | Qwen/Qwen3-0.6B | RTX 5060 Laptop 8151 MiB

长度筛选后随机采样 10 条，seed=20261008，Poisson $\lambda=0.5$ 请求/秒。

输入保留 trace 长度；输出上限保留 trace output_length，实际输出允许提前停止。

10/10 成功；平均 latency 1.3982 秒；最大客户端发送偏差 4.38 ms。

未测量 TTFT、缓存命中和显存收益；不将该小实验推广到完整 trace。

ID	trace输入	trace输出/上限	服务输入	实际输出	status	latency(s)
1	1309	451	1317	332	ok	4.2207
2	902	370	910	286	ok	2.2993
3	1314	342	1322	342	ok	2.4694
4	2957	14	2965	14	ok	0.2117
5	926	138	934	138	ok	0.8104
6	2013	432	2021	221	ok	1.1936
7	899	26	907	26	ok	0.1343
8	1704	389	1712	260	ok	1.2858
9	1581	283	1589	232	ok	1.1345
10	890	8	898	8	ok	0.2226

验证：重新生成采样索引与到达计划，与保存数据一致；请求实际发生重叠。

服务输入与本地 chat template 计数逐条一致，均比正文多 8 token。

默认采样后端缺少 nvcc 的失败轮已独立保留；正式轮使用 PyTorch 采样。

结果依据：results.jsonl / config.json / environment.json / validation.json。